

# A Quick Introduction to Database design with E/R (extracurricular to MAS 201)

## Data Structure: Relational Model

- **Relational Databases:**  
Schema + Data
- **Schema:**
  - collection of *tables*  
(also called *relations*)
  - each table has a set  
of *attributes*
  - no repeating relation names,  
no repeating attributes in  
one table
- **Data** (also called *instance*):
  - set of *tuples*
  - tuples have one *value* for each  
attribute

Movie			
ID	Title	Director	Actor
1	Wild	Lynch	Winger
2	Sky	Berto	Winger
3	Reds	Beatty	Beatty
4	Tango	Berto	Brando
5	Tango	Berto	Winger
7	Tango	Berto	Snyder

Schedule		
ID	Theater	Movie
1	Odeon	1
2	Forum	3
3	Forum	2

## Data Structure: Primary Keys; Foreign Keys are value-based pointers

Schedule			Movie			
ID	Theater	Movie	ID	Title	Director	Actor
1	Odeon	1	1	Wild	Lynch	Winger
2	Forum	3	2	Sky	Berto	Winger
3	Forum	2	3	Reds	Beatty	Beatty
			4	Tango	Berto	Brando
			5	Tango	Berto	Winger
			7	Tango	Berto	Snyder

- "ID is *primary key* of **Schedule**" => its value is unique in **Schedule.ID**
- "**Schedule.Movie** is foreign key (referring) to **Movie.ID**" means every **Movie** value of **Schedule** also appears as **Movie.ID**
- Intuitively, **Schedule.Movie** operates as pointer to **Movie(s)**

3

## Schema design has its own intricacies

Schedule			Movie			
ID	Theater	Movie	ID	Title	Director	Actor
1	Odeon	1	1	Wild	Lynch	Winger
2	Forum	3	2	Sky	Berto	Winger
3	Forum	2	3	Reds	Beatty	Beatty
			4	Tango	Berto	Brando
			5	Tango	Berto	Winger
			7	Tango	Berto	Snyder

- This is a bad schema design!
- Problems
  - Change the name of a theater
  - Change the name of a movie's director
  - What about theaters that play no movie?

4

## How to Design a Database and Avoid Bad Decisions

- With experience...
  - sweat, tears, etc etc
- Learning the normalization rules of database design
  - a well-developed mathematical theory about how to fix step by step a “bad” schema
- Think **entities and relationships** – then translate to relations
  - A practically useful way to come up with the good schema

5

## Data Structure: Relational Model

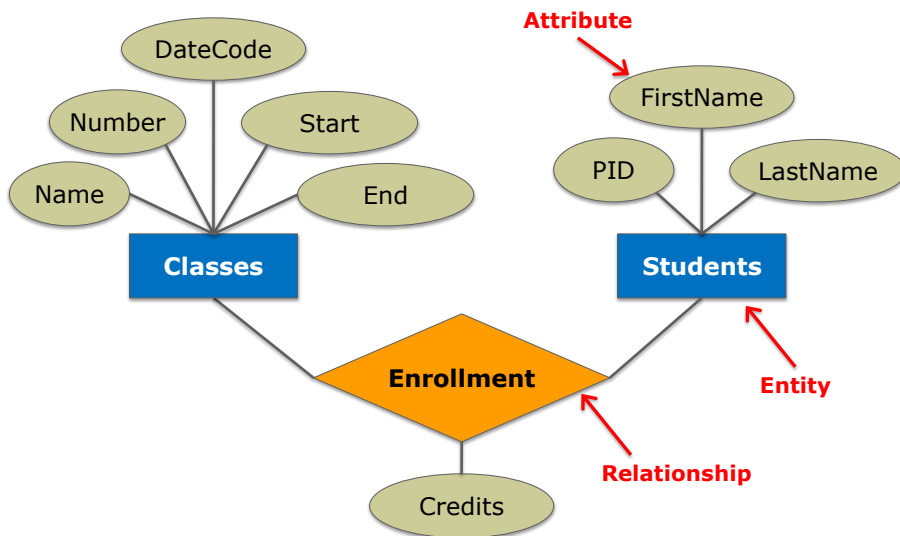
### Example Problem:

- Represent the students and Spring classes of the CSE department, including the enrollment of students in classes.
- Students have pid, first name and last name.
- Classes have a name, a number, date code (TR, MW, MWF) and start/end time.
- A student enrolls for a number of credits in a class.

### Solution:...

6

## Example 1: E/R-Based Design



7

## E/R → Relational Schema: Basic Translation

- For every entity
  - create corresponding table
  - For each attribute of the entity, add a corresponding attribute in the table
  - Include an ID attribute in the table even if not in E/R
- For every relationship
  - create corresponding table
  - For each attribute of the relationship, add a corresponding attribute in the table
  - For each referenced entity  $E_i$  include in the table a *required foreign key* attribute referencing ID of  $E_i$

8

## Sample relational database, per previous page's algorithm

Classes					
id	name	number	date_code	start_time	end_time
1	Web stuff	CSE135	TuTh	2:00	3:20
2	Databases	CSE132A	TuTh	3:30	4:50
4	VLSI	CSE121	F	<i>null</i>	<i>null</i>

Enrollment			
id	class	student	credits
1	1	1	4
2	1	2	3
3	4	3	4
4	1	3	3

Students			
id	pid	first_name	last_name
1	8888888	John	Smith
2	1111111	Mary	Doe
3	2222222	<i>null</i>	Chen

9

## Declaration of schemas in SQL's Data Definition Language

```

CREATE TABLE classes (
    ID          SERIAL PRIMARY KEY,
    name        TEXT,
    number      TEXT,
    date_code   TEXT,
    start_time  TIME,
    end_time    TIME
)
CREATE TABLE students (
    ID          SERIAL PRIMARY KEY,
    pid         TEXT,
    first_name  TEXT,
    last_name   TEXT
)
CREATE TABLE enrollment (
    ID          SERIAL,
    class       INTEGER REFERENCES classes (ID) NOT NULL,
    student     INTEGER REFERENCES students (ID) NOT NULL,
    credits     INTEGER
)

```

If we had "ID **INTEGER PRIMARY KEY**" we would be responsible for coming up with ID values. **SERIAL** leads to a counter that automatically provides ID values upon insertion of new tuples

Changed name from "end" to "end\_time" since "end" is reserved keyword

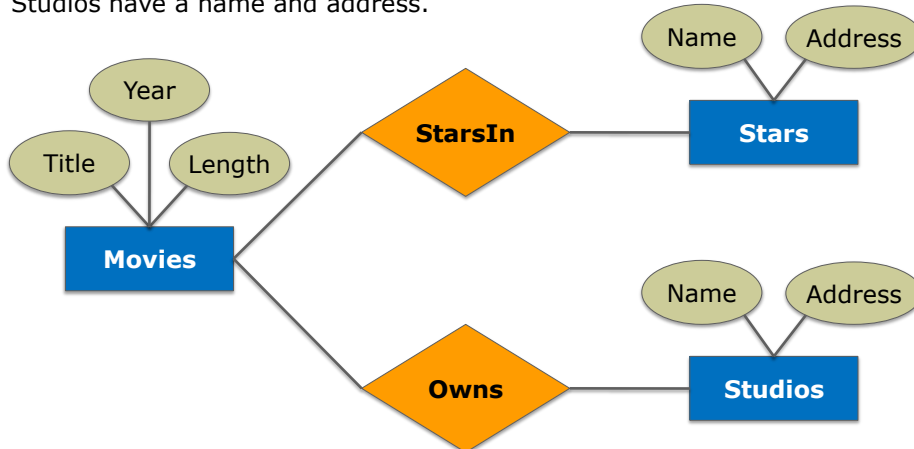
Foreign key declaration: Every value of **enrollment.class** must also appear as **classes.ID**

Declaration of "required" constraint: **enrollment.student** cannot be null (notice, it would make no sense to have an enrollment tuple without a student involved)

10

## Example 2a

Movies have a title, a year of release and length (in minutes).  
Actors have names and address.  
Actors appear in movies.  
A movie is (co-)owned by studios.  
Studios have a name and address.



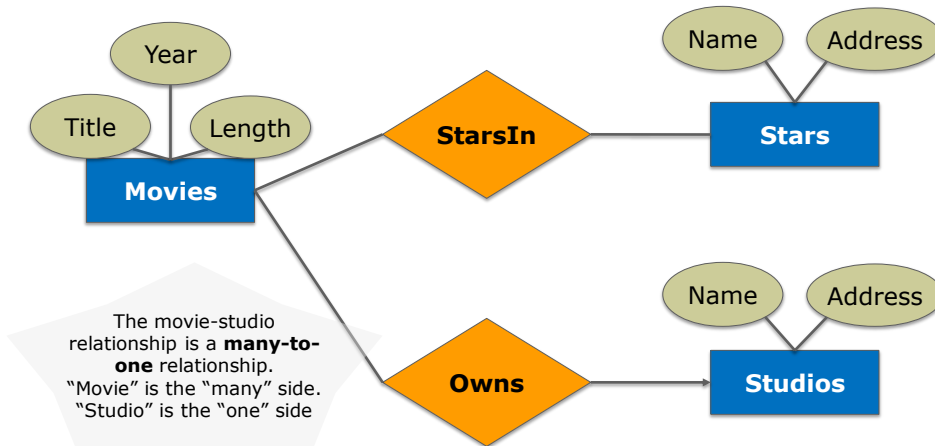
11

```
CREATE TABLE movies (
  ID          SERIAL PRIMARY KEY,
  title       TEXT,
  year        INTEGER,
  length      INTEGER,
)
CREATE TABLE stars (
  ID          SERIAL PRIMARY KEY,
  name        TEXT,
  address     TEXT
)
CREATE TABLE studios (
  ID          SERIAL PRIMARY KEY,
  name        TEXT,
  address     TEXT
)
CREATE TABLE starsin (
  ID          SERIAL,
  movie       INTEGER REFERENCES movies (ID) NOT NULL,
  star        INTEGER REFERENCES stars (ID) NOT NULL
)
CREATE TABLE ownership (
  ID          SERIAL,
  movie       INTEGER REFERENCES movies (ID) NOT NULL,
  owner       INTEGER REFERENCES studios (ID) NOT NULL
)
```

12

## Example 2b: many-to-one relationship

Modification to Example 2a:  
A movie is owned by **at most one** studio.



13

## E/R → Relational: Basic Translation revisited for many-to-one relationship

- For every entity, do the usual...
- For every **many-to-many** relationship, do the usual...
- For every **2-way many-to-one** relationship, where
  - $E_m$  is the "many" side
  - $E_o$  is the "one" side (pointed by the arrow)
  - **do not** create table, instead:
  - In the table corresponding to  $E_m$  add a (non-required) foreign key attribute referencing the ID of the table corresponding to  $E_o$

14

```

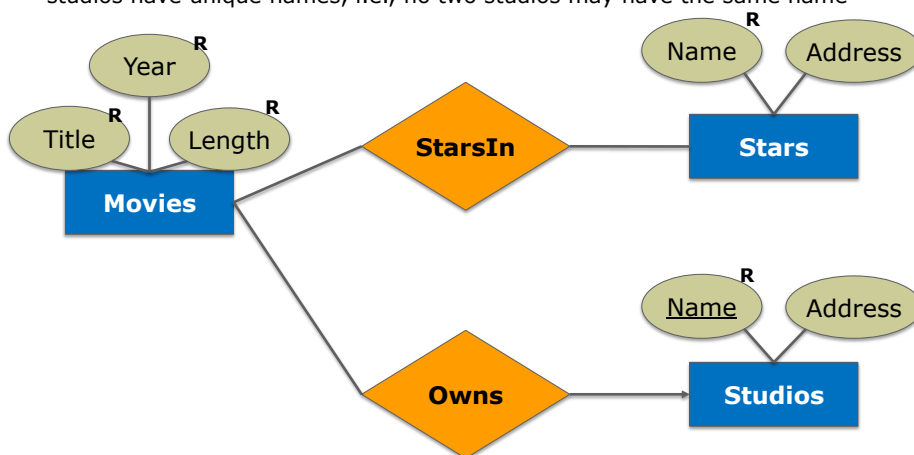
CREATE TABLE movies (
    ID          SERIAL PRIMARY KEY,
    title       TEXT,
    year        INTEGER,
    length      INTEGER,
    owner       INTEGER REFERENCES studios (ID)
)
CREATE TABLE stars (
    ID          SERIAL PRIMARY KEY,
    name        TEXT,
    address     TEXT
)
CREATE TABLE studios (
    ID          SERIAL PRIMARY KEY,
    name        TEXT,
    address     TEXT
)
CREATE TABLE starsin (
    ID          SERIAL,
    movie       INTEGER REFERENCES movies (ID) NOT NULL,
    star        INTEGER REFERENCES stars (ID) NOT NULL
)

```

15

## Example 2c: Constraints: uniqueness; required attributes

- In addition to Example 2b's assumptions, let us also assume that:
- title, year, length, star name and studio name are required attributes of the respective entities
    - default is that an attribute value may be **null**
  - studios have unique names, i.e., no two studios may have the same name




16



## Example 2d: Constraints: Required relationship; cardinality ranges

In addition to Example 2c's assumptions, let us also assume that:

- a movie is owned by **exactly one** studio
  - so far we had not assumed that the owning studio has to be known (not null)
- a movie must have at least one actor and no more than 100



17

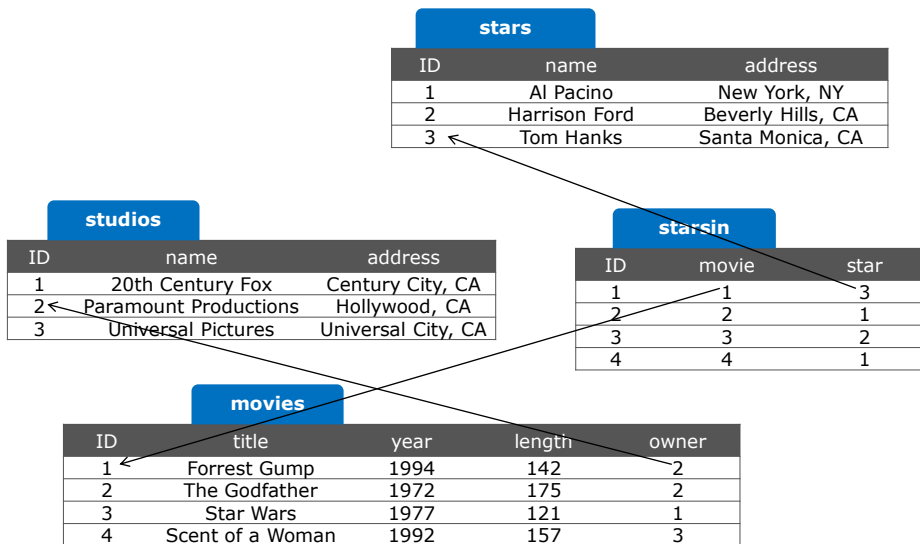
## SQL Schema for Examples 2c, 2d

```

CREATE TABLE movies (
    ID          SERIAL PRIMARY KEY,
    title       TEXT NOT NULL,
    year        INTEGER NOT NULL,
    length      INTEGER NOT NULL,
    owner       INTEGER REFERENCES studios (ID) NOT NULL
)
CREATE TABLE stars (
    ID          SERIAL PRIMARY KEY,
    name        TEXT NOT NULL,
    address     TEXT
)
CREATE TABLE studios (
    ID          SERIAL PRIMARY KEY,
    name        TEXT NOT NULL UNIQUE,
    address     TEXT
)
CREATE TABLE starsin (
    ID          SERIAL,
    movie       INTEGER REFERENCES movies (ID) NOT NULL,
    star        INTEGER REFERENCES stars (ID) NOT NULL
)
    
```

18

## A sample database



19

## Why do we want constraints? What happens when they are violated?

- Protect the database from erroneous data entry
- Prevent database states that are inconsistent with the rules of the business process you want to capture
- Whenever you attempt to change (insert, delete, update) the database in a way that violates a constraint the database will prevent the change
  - Try it out on the sample databases of the class page

20

## **Some constraints are not implemented by some SQL database systems**

- Consider the cardinality constraint that a movie has between 1 and 100 actors.
- The SQL standard provides a way, named CHECK constraints, to declare such
  - its specifics will make more sense once we have seen SQL queries
- However, no open source database implements the CHECK constraints.

21

## **Vice versa: SQL allows some constraints that are not in plain E/R**

Notable cases:

- Attribute value ranges
  - Example: Declare that the year of movies is after 1900
- Multi-attribute UNIQUE
  - Example: Declare that the (title, year) attribute value combination is unique

22

## Added constraints of previous slide to schema of Example 2d

```
CREATE TABLE movies (  
  ID          SERIAL PRIMARY KEY,  
  title       TEXT NOT NULL,  
  year        INTEGER NOT NULL CHECK (year > 1900),  
  length      INTEGER NOT NULL,  
  owner       INTEGER REFERENCES studios (ID) NOT NULL,  
  UNIQUE (title, year)  
)  
  
CREATE TABLE stars (  
  ID          SERIAL PRIMARY KEY,  
  name        TEXT NOT NULL,  
  address     TEXT  
)  
  
CREATE TABLE studios (  
  ID          SERIAL PRIMARY KEY,  
  name        TEXT NOT NULL UNIQUE,  
  address     TEXT  
)  
  
CREATE TABLE starsin (  
  ID          SERIAL,  
  movie       INTEGER REFERENCES movies (ID) NOT NULL,  
  star        INTEGER REFERENCES stars (ID) NOT NULL  
)
```

23

## Example 3: one-to-one relationships

Assume that a president manages exactly one studio and  
a studio may have at most one president.  
Notice: a studio may not have a president but  
in order to be a president one has to manage a studio.



24

## 1<sup>st</sup> candidate

```
CREATE TABLE presidents (  
  ID          SERIAL PRIMARY KEY,  
  name        TEXT,  
  age         INTEGER,  
  manages     INTEGER REFERENCES studios (ID) NOT NULL UNIQUE  
)  
  
CREATE TABLE studios (  
  ID          SERIAL PRIMARY KEY,  
  name        TEXT,  
  address     TEXT  
)
```

Guarantees that  
in order to be  
president, one  
has to manage a  
studio

Guarantees that  
no two presidents  
may manage the  
same studio

25


## 2<sup>nd</sup> candidate

2<sup>nd</sup> candidate is not preferred. Why? What constraint it misses?

```
CREATE TABLE presidents (  
  ID          SERIAL PRIMARY KEY,  
  name        TEXT,  
  age         INTEGER  
)  
  
CREATE TABLE studios (  
  ID          SERIAL PRIMARY KEY,  
  name        TEXT,  
  address     TEXT,  
  managedBy   INTEGER REFERENCES presidents (ID) UNIQUE  
)
```

26

## Example 4: 3-Way Relationship



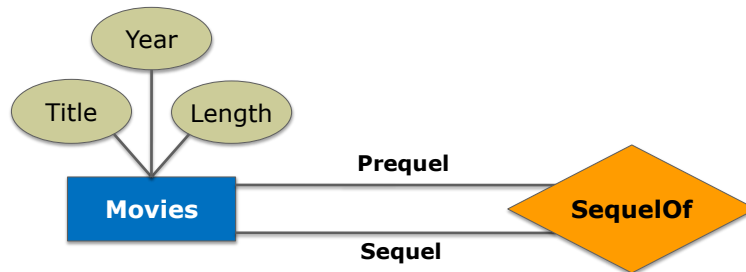
- A studio has contracted with a particular star to act in a particular movie
  - No ownership of movies by studios

27

```
CREATE TABLE contract (  
  ID          SERIAL,  
  movie       INTEGER REFERENCES movies (ID) NOT NULL,  
  star        INTEGER REFERENCES stars (ID) NOT NULL,  
  owner       INTEGER REFERENCES studios (ID) NOT NULL,  
  fee         INTEGER  
)
```

28

## Example 5a : Self-Relationships with Roles



29


```
CREATE TABLE movies (
    ID          SERIAL PRIMARY KEY,
    ...
)

CREATE TABLE sequelof (
    ID          SERIAL,
    prequel     INTEGER REFERENCES movies (ID) NOT NULL,
    sequel      INTEGER REFERENCES movies (ID) NOT NULL
)
```

Notice the use of  
roles as attributes  
names for the  
foreign keys

30

## Example 5b : Combo: One-to-one Self-Relationship



A movie has at most one direct "prequel" and at most one direct "sequel"

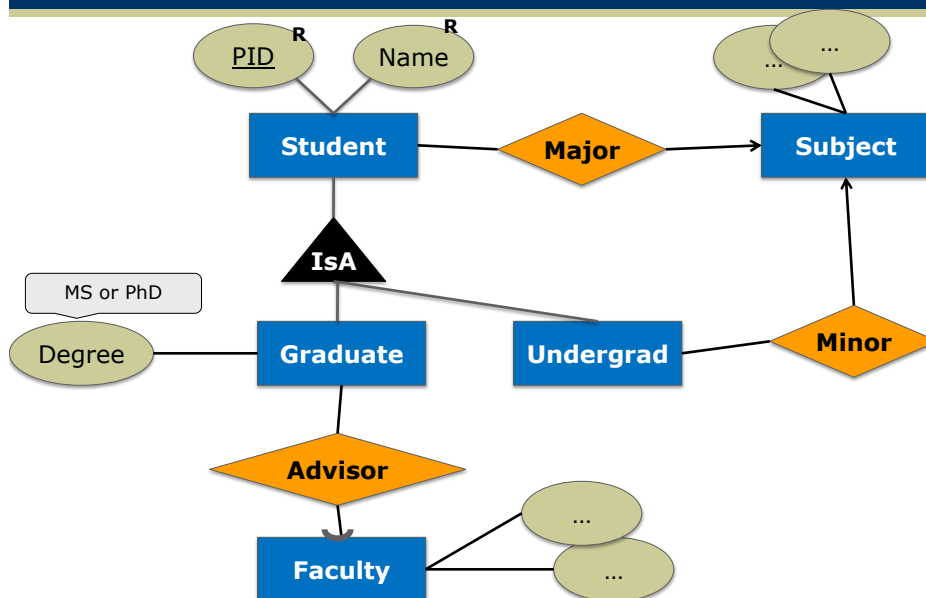
Modeling movie sequels by "DirectSequelOf" is preferable to using "SequelOf" of previous slide

A lesson about database design:

- Good designs avoid redundancy.
- No stored piece of data should be inferable from other stored pieces of data

31

## Example 6: Subclassing



32



## Schemas for subclassing: Candidate 1

```
CREATE TABLE student(
    ID        SERIAL PRIMARY KEY,
    pid       TEXT NOT NULL UNIQUE,
    name      TEXT NOT NULL,
    major     INTEGER REFERENCES subject(ID)
)
CREATE TABLE undergrad(
    studentid  INTEGER REFERENCES student(ID) NOT NULL,
    minor      INTEGER REFERENCES subject(ID)
)
CREATE TABLE graduate(
    studentid  INTEGER REFERENCES student(ID) NOT NULL,
    degree    TEXT NOT NULL CHECK (degree="PhD" OR degree="MS"),
    advisor    INTEGER REFERENCES faculty(ID) NOT NULL
)
CREATE TABLE subject(
    ID        SERIAL PRIMARY KEY,
    ...
)
CREATE TABLE faculty(
    ID        SERIAL PRIMARY KEY,
    ...
)
```

+ captures constraints  
- Information about graduates is spread on two tables  
- Creating a report about students is a tricky query  
To appreciate the above wait till we discuss SQL

33

## Schemas for subclassing: Candidate 2

```
CREATE TABLE student(
    ID        SERIAL PRIMARY KEY,
    pid       TEXT NOT NULL UNIQUE,
    name      TEXT NOT NULL,
    kind      CHAR(1) CHECK (kind='U' OR kind='S'),
    major     INTEGER REFERENCES subject(ID),
    minor     INTEGER REFERENCES subject(ID),
    degree    TEXT CHECK (degree="PhD" OR degree="MS"),
    advisor    INTEGER REFERENCES faculty(ID)
)
CREATE TABLE subject(
    ID        SERIAL PRIMARY KEY,
    ...
)
CREATE TABLE faculty(
    ID        SERIAL PRIMARY KEY,
    ...
)
```

-misses constraints  
E.g., notice that it does not capture that a graduate student must have an advisor since we had to make the advisor attribute non-required in order to accommodate having undergraduates in the same table

34

## Not covered E/R features

- Weak entities
  - double-lined entities and relationships
- Necessary participation of entity in relationship
- ... more