

A quick guide to writing SQL queries

MAS 201

Access (Query) & Modification Language: SQL

- SQL
 - used by the database user
 - **declarative**: we only describe **what** we want to retrieve
 - based on tuple relational calculus
- The result of a query is a table (regardless of the query language used)

SQL Queries: Basic One-table

- Basic form

SELECT A_1, \dots, A_N

FROM R

WHERE <condition>

- **WHERE** clause is optional
- Find all tuples of R that satisfy the (boolean) condition and return their attributes A_1, \dots, A_N

Find first names and last names of all students

SELECT first_name, last_name
FROM students;

Find all students whose first name is John; project all attributes

SELECT *
FROM students
WHERE first_name = 'John';

3

SQL Queries: Putting together multiple tables

- Basic form

SELECT A_1, \dots, A_N

FROM R_1, \dots, R_M

WHERE <condition>

- When more than one relations in the **FROM** clause have an attribute named A , we refer to a specific A attribute as <RelationName>. A
- Hardest to get used to, yet most important feature of SQL

Produce a table that shows the pid, first name and last name of every student enrolled in the class with ID 1, along with the number of credit units in the "class 1" enrollment

4

(repeat)

Classes						
id	name	number	date_code	start_time	end_time	
1	Web stuff	CSE135	TuTh	2:00	3:20	
2	Databases	CSE132A	TuTh	3:30	4:50	
4	VLSI	CSE121	F	null	null	

Enrollment			
id	class	student	credits
1	1	1	4
2	1	2	3
3	4	3	4
4	1	3	3

Students			
id	pid	first_name	last_name
1	8888888	John	Smith
2	1111111	Mary	Doe
3	2222222	null	Chen

5

Take One: Understanding FROM as producing all combinations

```

SELECT students.pid, students.first_name,
       students.last_name, enrollment.credits
FROM   students, enrollment
WHERE  students.id = enrollment.student
       AND enrollment.class = 1 ;

```

"FROM" produces all 12 tuples made from one "students" tuple and one "enrollment" tuple

Student part of the tuple				Enrollment part of the tuple			
Students. id	pid	first_name	last_name	Enrollment. id	class	student	credits
1	88..	John	Smith	1	1	1	4
1	88..	John	Smith	2	1	2	3
1	88..	John	Smith	3	4	3	4
1	88..	John	Smith	4	1	3	3
2	11..	Mary	Doe	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
2	11..	Mary	Doe	3	4	3	4
2	11..	Mary	Doe	4	1	3	3
3	22..	null	Chen	1	1	1	4
3	22..	null	Chen	2	1	2	3
3	22..	null	Chen	3	4	3	4
3	22..	null	Chen	4	1	3	3

6

Understanding WHERE as qualifying the tuples that satisfy the condition

Students. id	pid	first_name	last_name	Enrollment. id	class	student	credits
1	88..	John	Smith	1	1	1	4
1	88..	John	Smith	2	1	2	3
1	88..	John	Smith	3	4	3	4
1	88..	John	Smith	4	1	3	3
2	11..	Mary	Doe	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
2	11..	Mary	Doe	3	4	3	4
2	11..	Mary	Doe	4	1	3	3
3	22..	null	Chen	1	1	1	4
3	22..	null	Chen	2	1	2	3
3	22..	null	Chen	3	4	3	4
3	22..	null	Chen	4	1	3	3

7

Understanding SELECT as keeping the listed columns (highlighted below)

Students. id	pid	first_name	last_name	Enrollment. id	class	student	credits
1	88..	John	Smith	1	1	1	4
1	88..	John	Smith	2	1	2	3
1	88..	John	Smith	3	4	3	4
1	88..	John	Smith	4	1	3	3
2	11..	Mary	Doe	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
2	11..	Mary	Doe	3	4	3	4
2	11..	Mary	Doe	4	1	3	3
3	22..	null	Chen	1	1	1	4
3	22..	null	Chen	2	1	2	3
3	22..	null	Chen	3	4	3	4
3	22..	null	Chen	4	1	3	3

Net result of the query is

Students .pid	Students.first_name	Students.last_name	Enrollment.credits
88..	John	Smith	4
11..	Mary	Doe	3
22..	null	Chen	3

8

Generalize to any number of tables

Produce a table that shows the pid, first name and last name of every student enrolled in the CSE135 class along with the number of credit units in his/her enrollment

HOW TO UNDERSTAND THE FROM AND WHERE (AT LEAST UNTIL WE TALK ABOUT DUPLICATES):

Find the students, whose students.id appears in an enrollment tuple as enrollment.student, and the enrollment.class of this tuple is the class.id of a class tuple whose number is CSE135

```
SELECT students.pid, students.first_name,
       students.last_name, enrollment.credits
FROM   students, enrollment, classes
WHERE  classes.number = 'CSE135'
       AND students.id = enrollment.student
       AND enrollment.class = classes.id ;
```

9

You can omit table names in SELECT, WHERE when attribute is unambiguous

```
SELECT pid, first_name, last_name, credits
FROM   students, enrollment, classes
WHERE  number = 'CSE135'
       AND students.id = student
       AND class = classes.id ;
```

10

Take Two on the previous exercises: The algebraic way to express joins

Produce a table that shows the pid, first name and last name of every student enrolled in the class with ID 1, along with the number of credit units in the "class 1" enrollment

```
SELECT students.pid, students.first_name,
       students.last_name, enrollment.credits
FROM   students JOIN enrollment
       ON students.id = enrollment.student
WHERE  enrollment.class = 1 ;
```

11

Take two cont'd

FROM clause result

Student part of the tuple				Enrollment part of the tuple			
Students. id	pid	first_name	last_name	Enrollment. id	class	student	credits
1	88..	John	Smith	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
3	22..	null	Chen	3	4	3	4
3	22..	null	Chen	4	1	3	3

SELECT clause result

Students. id	pid	first_name	last_name	Enrollment. id	class	student	credits
1	88..	John	Smith	1	1	1	4
2	11..	Mary	Doe	2	1	2	3
3	22..	null	Chen	3	4	3	4
3	22..	null	Chen	4	1	3	3

Net result of the query is

Students. pid	Students.first_name	Students.last_name	Enrollment.credits
88..	John	Smith	4
11..	Mary	Doe	3
22..	null	Chen	3

12

Take two (algebraic approach)

Second example

Produce a table that shows the pid, first name and last name of every student enrolled in the CSE135 class along with the number of credit units in his/her 135 enrollment

Take One:

```
SELECT students.pid, students.first_name,  
       students.last_name, enrollment.credits  
FROM   students, enrollment, classes  
WHERE  classes.number = 'CSE135'  
       AND students.id = enrollment.student  
       AND enrollment.class = classes.id
```

Take Two:

```
SELECT students.pid, students.first_name,  
       students.last_name, enrollment.credits  
FROM   (students JOIN enrollment ON student.id = enrollment.student)  
       JOIN classes ON enrollment.class = class.id  
WHERE  classes.number = 'CSE135'
```

13

Heuristics on writing queries

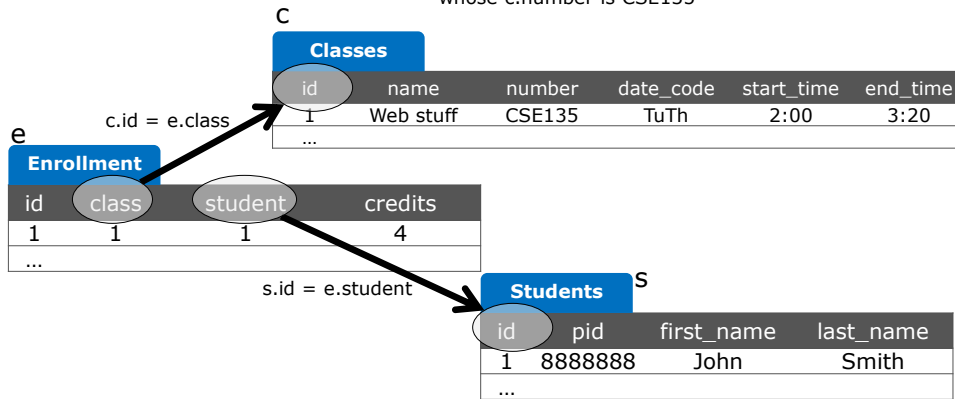
- Have you reached the point where you understand how queries work but have difficulty writing queries yourself?
- The following heuristics will help you translate a requirement expressed in English into a query
 - The key point is to translate informal English into a precise English statement about what tuples you expect to find in the database

14

Hints for writing FROM/WHERE: Rephrase the statement, see it as a navigation across primary/foreign keys

Produce a table that shows the pid, first name and last name of every student enrolled in the CSE135 class along with the number of credit units in his/her 135 enrollment

- Find any students tuple s,
- that is connected to an enrollment tuple e
 - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
 - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is CSE135



15

- Find any students tuple s, FROM students AS s
- that is connected to an enrollment tuple e
 - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
 - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is CSE135

- Find any students tuple s,
- that is connected to an enrollment tuple e
 - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
 - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is CSE135

Take One: Declarative

FROM students AS s,
enrollment AS e
WHERE s.id = e.student

Take Two: Algebraic

FROM students AS s
JOIN enrollment AS e
ON s.id = e.student

16

- Find any students tuple s,
- that is connected to an enrollment tuple e
 - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
 - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is CSE135

Take One: Declarative

```
FROM students AS s,
      enrollment AS e,
      classes AS c
WHERE s.id = e.student
      AND c.id = e.class
```

Take Two: Algebraic

```
FROM ( students AS s
      JOIN enrollment AS e
        ON s.id = e.student )
      JOIN classes AS c
        ON c.id = e.class
```

17

- Find any students tuple s,
- that is connected to an enrollment tuple e
 - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
 - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is CSE135

Take One: Declarative

```
FROM students AS s,
      enrollment AS e,
      classes AS c
WHERE s.id = e.student
      AND c.id = e.class
      AND c.number = 'CSE135'
```

Take Two: Algebraic

```
FROM ( students AS s
      JOIN enrollment AS e
        ON s.id = e.student )
      JOIN classes AS c
        ON c.id = e.class
WHERE c.number = 'CSE135'
```

18

- Find any students tuple s, FROM students AS s
- that is connected to an enrollment tuple e
 - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
 - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is CSE135

- Find any students tuple s,
- that is connected to an enrollment tuple e
 - i.e., whose s.id appears in an enrollment tuple e as e.student,
- and e is connected to a classes tuple c
 - i.e., the e.class of e appears as c.id of the tuple c,
- whose c.number is CSE135

Take One: Declarative

```
FROM students AS s,
      enrollment AS e
WHERE s.id = e.student
```

Take Two: Algebraic

```
FROM students AS s,
      JOIN enrollment AS e
ON s.id = e.student
```

19

SQL Queries: Nesting

- The **WHERE** clause can contain predicates of the form
 - **attr/value IN <query>**
 - **attr/value NOT IN <query>**
- The predicate is satisfied if the **attr** or **value** appears in the result of the nested **<query>**
- Also
 - **EXISTS <query>**
 - **NOT EXISTS <query>**

20

Nested subquery example (uncorrelated subquery)

Produce a table that shows the pid, first and last name of every student enrolled in CSE135

```
SELECT pid, first_name, last_name
FROM students
WHERE id IN
  ( SELECT student
    FROM enrollment, classes
    WHERE number='CSE135'
      AND class=classes.id
  )
```

"Uncorrelated" in the sense that the nested query could be a standalone query

Nested queries help modularize the task:
Nested query finds the id's of the students who take CSE135. Then the outer query prints out pid and name for every student whose id appears in the result of the nested query

21

Nested subquery example, correlated

```
SELECT pid, first_name, last_name
FROM students
WHERE EXISTS
  ( SELECT *
    FROM enrollment, classes
    WHERE number='CSE135'
      AND class=classes.id
      AND student = students.id
  )
```

Correlation of nested query to outside query. The nested query is not a standalone.

There may be IN queries that are correlated and EXISTS queries that are uncorrelated.

22

SQL Queries, advanced: Aliases

- Use the same relation more than once in the **FROM** clause
- Tuple variables
- **Problem:** Find the other classes taken by students who take CSE135
 - First, also showing the students, i.e., produce a table where each row has the name of a 135 student and the name of another class he/she takes

23

produce a table where each row has the name of a 135 student and the name of another class he/she takes

```
SELECT c_others.name, first_name, last_name
FROM   classes AS c_135, enrollment AS e_135,
       students,
       enrollment AS e_others, classes AS c_others
WHERE  c_135.number = 'CSE135'
       AND c_135.id = e_135.class
       AND e_135.student = students.id
       AND students.id = e_others.student
       AND e_others.class = c_others.id
       AND NOT (c_others.number = 'CSE135')
```

24

Second, show just the other classes. Notice use of **DISTINCT**

```
SELECT DISTINCT c_others.name
FROM   classes AS c_135, enrollment AS e_135,
       enrollment AS e_others, classes AS c_others
WHERE  c_135.number = 'CSE135'
       AND c_135.id = e_135.class
       AND e_135.student = e_others.student
       AND e_others.class = c_others.id
       AND NOT (c_others.number = 'CSE135')
```

25

Use of nested subqueries may reduce need for aliases => easier to write, read

Find the CSE135 students who take a Friday 11:00 am class

```
SELECT first_name, last_name
FROM   students, enrollment, classes
WHERE  students.id = student
       AND class = classes.id
       AND number = 'CSE135'
       AND students.id IN
       (
         SELECT student
         FROM   enrollment, classes
         WHERE  classes.id = class
               AND date_code = 'F'
               AND start_time = '11:00'
       )
```

Nested query
computes the id's of
students enrolled in
Friday 11:00AM classes

26

SQL Queries: Aggregation & Grouping

- Aggregate functions: SUM, AVG, COUNT, MIN, MAX, and recently user defined functions as well
- GROUP BY

Employee		
Name	Dept	Salary
Joe	Toys	45
Nick	PCs	50
Jim	Toys	35
Jack	PCs	40

Example: Find the average salary of all employees:

```
SELECT AVG(Salary) AS AvgSal  
FROM Employee
```

AvgSal
42.5

Example: Find the average salary for each department:

```
SELECT Dept, AVG(Salary) AS AvgSal  
FROM Employee  
GROUP BY Dept
```

Dept	AvgSal
Toys	40
PCs	45

27

SQL Grouping: Conditions that Apply on Groups

- HAVING <condition> may follow a GROUP BY clause
- If so, the condition applies to each group, and groups not satisfying the condition are eliminated
- Example:** Find the average salary in each department that has more than 1 employee:

```
SELECT Dept, AVG(Salary) AS AvgSal  
FROM Employee  
GROUP BY Dept  
HAVING COUNT(Name) >1
```

28

Let's mix features we've seen: Aggregation after joining tables

- **Problem:** List all enrolled students and the number of total credits for which they have registered

```
SELECT  students.id, first_name, last_name, SUM(credits)
FROM    students, enrollment
WHERE   students.id = enrollment.student
GROUP BY students.id, first_name, last_name
```

29

The outerjoin operator

- New construct in FROM clause
- R LEFT OUTER JOIN S ON R.<attr of R>=S.<attr of J>
- R FULL OUTER JOIN S ON R.<attr of R>=S.<attr of J>

R		S	
RJ	RV	SJ	SV
1	RV1	1	SV1
2	RV2	3	SV3

SELECT *			
FROM R LEFT OUTERJOIN S ON R.RJ=S.SJ			
RJ	RV	SJ	SV
1	RV1	1	SV1
2	RV2	Null	Null

SELECT *			
FROM R FULL OUTERJOIN S ON R.RJ=S.SJ			
RJ	RV	SJ	SV
1	RV1	1	SV1
2	RV2	Null	Null
Null	Null	3	SV3

30

An application of outerjoin

- **Problem:** List all students and the number of total credits for which they have registered
 - Notice that you must also list non-enrolled students

```
SELECT  students.id, first_name, last_name, SUM(credits)
FROM    students LEFT OUTER JOIN enrollment ON
        students.id = enrollment.student
GROUP BY students.id, first_name, last_name
```

31

SQL: More Bells and Whistles ...

- Pattern matching conditions
 - <attr> LIKE <pattern>

Retrieve all students whose name contains "Sm"

```
SELECT *
FROM Students
WHERE name LIKE '%Sm%'
```

32

...and a Few “Dirty” Points

- **Null values**

- All comparisons involving NULL are **false** by definition
- All aggregation operations, except `COUNT(*)`, ignore NULL values

33

Null Values and Aggregates

- Example:

R	
a	b
x	1
x	2
x	null
null	null
null	null

```
SELECT COUNT(a) , COUNT(b) , AVG(b) , COUNT(*)  
FROM R  
GROUP BY a
```

count(a)	count(b)	avg(b)	count(*)
3	2	1.5	3
0	0	null	2

34

Universal Quantification by Negation (difficult)

Problem:

- Find the students that take **every** class 'John Smith' takes

Rephrase:

- Find the students such that there is no class that 'John Smith' takes and they do not take

35

Discussed in class and discussion section

How to solve in easy steps the following complex query:

Create a table that shows all time slots (date, start time, end time)
when students of CSE135 attend a lecture of another class;
Show also how many students attend a class at each time slot.

36

SQL as a Data Manipulation Language: Insertions

- Inserting tuples

```
INSERT INTO R (A1, ..., Ak)  
VALUES (v1, ..., vk);
```

- Some values may be left NULL

- Use results of queries for insertion

```
INSERT INTO R  
SELECT ...  
FROM ...  
WHERE ...
```

- Insert in Students 'John Doe' with A# 99999999

```
INSERT INTO students  
(pid, first_name, last_name)  
VALUES  
(99999999, 'John', 'Doe')
```

- Enroll all CSE135 students into CSE132A

```
INSERT INTO enrollment (class,  
student)  
SELECT c132a.id, student  
FROM classes AS c135, enrollment,  
classes AS c132a  
WHERE c135.number='CSE135' AND  
enrollment.class=c135.id AND  
c132a.number='CSE132A'
```

37

SQL as a Data Manipulation Language: Updates and Deletions

- Deletion basic form: delete every tuple that satisfies <cond>:

```
DELETE FROM R  
WHERE <cond>
```

- Update basic form: update every tuple that satisfies <cond> in the way specified by the SET clause:

```
UPDATE R  
SET A1=<exp1>, ..., Ak=<expk>  
WHERE <cond>
```

- Delete "John" "Smith"
- DELETE FROM students WHERE first_name='John' AND last_name='Smith'

- Update the registered credits of all CSE135 students to 5

```
UPDATE enrollment  
SET credits=5  
WHERE class=1
```

```
UPDATE enrollment  
SET credits=5  
WHERE class IN  
(SELECT id FROM classes  
WHERE number='CSE135')
```

38